

Quiz 3

GLM Theory

```
library(tidyverse)
```

```
-- Attaching core tidyverse packages ----- tidyverse 2.0.0 --
v dplyr      1.1.4      v readr      2.1.5
v forcats    1.0.0      v stringr    1.5.1
v ggplot2    3.5.1      v tibble     3.2.1
v lubridate  1.9.3      v tidyr      1.3.1
v purrr      1.0.2
-- Conflicts ----- tidyverse_conflicts() --
x dplyr::filter() masks stats::filter()
x dplyr::lag()     masks stats::lag()
i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to become
```

Question 1

Is it acceptable to use the quasi-poisson model when the data actually follow the Poisson distribution? Briefly explain why or why not.

It is not acceptable to use a quasi-poisson model when the data actually follows a Poisson distribution. In fact, the quasi-poisson exists for the case when the data doesn't satisfy assumptions of Poisson data, specifically the assumption that the mean and variance are equal and quasi-poisson will inflate the standard errors to handle variance greater than we would expect in Poisson data. If our data follows a Poisson distribution, then the mean=variance assumption should be satisfied, and quasi-poisson just takes away from the theoretical strength of our analysis.

Question 2

The support of the distribution is $y > 0$ or equivalently $(0, \infty)$.

Question 3

We start of with the PMF of the truncated Poisson, which is

$$P(Y = y) = \frac{e^{-\lambda} \lambda^y}{(1 - e^{-\lambda}) y!}$$

Note that our support is not dependent on an unknown parameter, meaning we can attempt to write this in exponential family form to see if the truncated Poisson is a member of the exponential family.

We will exponentiate the logarithm of the PMF.

$$\begin{aligned} & \exp\left(\log \frac{e^{-\lambda} \lambda^y}{(1 - e^{-\lambda}) y!}\right) \\ & \exp(\log(e^{-\lambda}) + \log(\lambda^y) - \log(1 - e^{-\lambda}) - \log(y!)) \end{aligned}$$

$$\frac{1}{y!} \exp\left(\log\left(\frac{e^{-\lambda}}{1 - e^{-\lambda}}\right) + y \log(\lambda)\right)$$

We notice that this belongs to the exponential family.

The exponential family in canonical form can be written as:

$$f(y \mid \lambda) = h(y) \exp(\eta(\lambda) T(y) - \psi(\lambda))$$

where

$$h(y) = \frac{1}{y!}$$

$$\psi(\lambda) = -\log\left(\frac{e^{-\lambda}}{1 - e^{-\lambda}}\right)$$

$$\eta = \log(\lambda)$$

$$T(y) = y$$

Question 4

Recall that the log partition function can be derived from the η and ψ components of our exponential family structure, specifically reparameterizing ψ in terms of η .

We found that $\eta = \log(\lambda)$ or equivalently that $e^\eta = \lambda$.

From $\psi = -\log(\frac{e^{-\lambda}}{1-e^{-\lambda}})$, we write the log partition function:

$$A = -\log\left(\frac{e^{-e^\eta}}{1 - e^{-e^\eta}}\right)$$

By logarithm rules, we can say

$$A = e^\eta + \log(1 - e^{-e^\eta})$$

From the log partition function, we can weaponize the following fact to find $E(Y)$

$$E(Y) = \frac{\partial A}{\partial \eta}$$

$$E(Y) = e^\eta + \frac{e^{-e^\eta} e^\eta}{1 - e^{-e^\eta}}$$

This is equivalent to

$$E(Y) = \frac{e^\eta}{1 - e^{-e^\eta}}$$

Since $e^\eta = \lambda$,

we can say

$$E(Y) = \frac{\lambda}{1 - e^{-\lambda}} = \frac{\lambda e^\lambda}{e^\lambda - 1}$$

Question 5

The canonical link can be viewed by η which connects our response to the predictors. Since $\eta = \log(\lambda)$, $X\beta = \log(\lambda)$ is the link function.

Question 6

To find the log likelihood of the truncated Poisson, I will first write out the likelihood by taking the product of PMF from $n=1, \dots, n$.

We define the PMF of the truncated Poisson as

$$f_Y = \frac{e^{-\lambda} \lambda^y}{(1 - e^{-\lambda}) y!}$$

The likelihood is then

$$L(\lambda) = \prod_{i=1}^n \frac{e^{-\lambda} \lambda^{y_i}}{(1 - e^{-\lambda}) y_i!}$$

$$L(\lambda) = \frac{e^{-n\lambda} \lambda^{\sum_{i=1}^n y_i}}{(1 - e^{-\lambda})^n \prod_{i=1}^n y_i!}$$

In attempt to find the log likelihood, we will take the logarithm.

$$\log L(\lambda) = \log\left(\frac{e^{-n\lambda} \lambda^{\sum_{i=1}^n y_i}}{(1 - e^{-\lambda})^n \prod_{i=1}^n y_i!}\right)$$

By logarithm properties

$$\log L(\lambda) = -n\lambda + \sum_{i=1}^n y_i \log \lambda - n \log(1 - e^{-\lambda}) - \sum_{i=1}^n \log(y_i!)$$

Using our link $\log(\lambda) = X\beta$, we know $\lambda = e^{X\beta}$, so we can equivalently write our log likelihood as

$$\log L(\beta) = -ne^{X\beta} + \sum_{i=1}^n y_i X\beta - n \log(1 - e^{-e^{X\beta}}) - \sum_{i=1}^n \log(y_i!)$$

This is equivalently

$$= \sum_{i=1}^n [-e^{X\beta} + y_i X\beta - \log(1 - e^{-e^{X\beta}}) - \log(y_i!)]$$

The vectorized form can be written as

$$\log L(\beta) = \sum_{i=1}^n [-e^{X_i^T \beta} + y_i X_i^T \beta - \log(1 - e^{-e^{X_i^T \beta}}) - \log(y_i!)]$$

This matches the stated log likelihood.

Question 7

To find the score function, I will take the partial derivative of the log likelihood with respect to β .

$$S(\beta) = \frac{\partial \log L(\beta)}{\partial \beta} = \sum_{i=1}^n [-e^{X_i^T \beta} X_i + y_i X_i - \frac{-e^{-e^{X_i^T \beta}} \times -e^{X_i^T \beta} X_i}{1 - e^{-e^{X_i^T \beta}}}]$$

$$S(\beta) = \frac{\partial \log L(\beta)}{\partial \beta} = \sum_{i=1}^n [-e^{X_i^T \beta} X_i + y_i X_i - \frac{e^{-e^{X_i^T \beta}} \times e^{X_i^T \beta} X_i}{1 - e^{-e^{X_i^T \beta}}}]$$

$$S(\beta) = \frac{\partial \log L(\beta)}{\partial \beta} = \sum_{i=1}^n [y_i X_i - \frac{e^{-e^{X_i^T \beta}} \times e^{X_i^T \beta} X_i + e^{X_i^T \beta} X_i (1 - e^{-e^{X_i^T \beta}})}{1 - e^{-e^{X_i^T \beta}}}]$$

This allows us to simplify our score function to

$$S(\beta) = \sum_{i=1}^n [y_i X_i - \frac{e^{X_i^T \beta}}{1 - e^{-e^{X_i^T \beta}}} X_i]$$

The score function in matrix form is:

$$S(\beta) = X^T (y - \frac{e^{X\beta}}{1 - e^{-e^{X\beta}}})$$

The hessian function is found by taking the second partial of the log likelihood with respect to β . This is also found by taking the partial derivative of the score function with respect to β

$$\begin{aligned} H(\beta) &= \frac{\partial}{\partial \beta} S(\beta) = \frac{\partial}{\partial \beta} [\sum_{i=1}^n [y_i X_i - \frac{e^{X_i^T \beta}}{1 - e^{-e^{X_i^T \beta}}} X_i]] \\ &= - \sum_{i=1}^n \frac{\partial}{\partial \beta} (\frac{e^{X_i^T \beta}}{1 - e^{-e^{X_i^T \beta}}} X_i) \end{aligned}$$

We can use quotient rule to find this derivative

$$\frac{\partial}{\partial \beta} (\frac{e^{X_i^T \beta}}{1 - e^{-e^{X_i^T \beta}}} X_i) = \frac{(1 - e^{-e^{X_i^T \beta}}) e^{X_i^T \beta} X_i X_i^T - e^{X_i^T \beta} e^{-e^{X_i^T \beta}} e^{X_i^T \beta} X_i}{(1 - e^{-e^{X_i^T \beta}})^2}$$

We use this to write the Hessian

$$H(\beta) = -\sum_{i=1}^n \frac{(e^{X_i^T \beta} - e^{-e^{X_i^T \beta}} e^{X_i^T \beta} - e^{2X_i^T \beta} e^{-e^{X_i^T \beta}})}{(1 - e^{-e^{X_i^T \beta}})^2} X_i X_i^T$$

In matrix form, we can write the Hessian as:

$$H(\beta) = -X^T W X$$

where $W = \frac{e^{X\beta} - e^{-e^{X\beta}} e^{X\beta} - e^{2X\beta} e^{-e^{X\beta}}}{(1 - e^{-e^{X\beta}})^2}$

Question 8

For simplicity, I use $m = \exp(X\beta)$ and $g = \frac{\exp(X\beta)}{(1 - \exp(-m))}$

```
trunc_pois_Newton_Raphson <- function(x, y, b_init = rep(1, ncol(x)),
                                     tol = 1*10^(-8))
{
  change <- Inf
  b_old <- b_init

  while(change > tol)
  {
    eta <- x %*% b_old
    em <- exp(eta)
    w <- diag(as.vector((em - (em * exp(-em)) - (em^2)*exp(-em))/(1 - exp(-em)^2))), nrow = nrow(x),
              ncol = nrow(x))
    g <- exp(eta)/(1 - exp(-em))

    b_new <- b_old + solve(t(x)%*% w %*% x) %*% t(x) %*% (y - g)

    change <- sqrt(sum((b_new - b_old)^2))

    b_old <- b_new
  }

  b_new
}
```

I will test my algorithm using the college dataset we used on a previous assignment.


```
college <- read_csv("data/college-data.csv")
```

```
Rows: 593 Columns: 15
```

```
-- Column specification -----
```

```
Delimiter: ","
```

```
chr (5): name, state, state_code, type, degree_length
```

```
dbl (10): room_and_board, in_state_tuition, in_state_total, out_of_state_tui...
```

```
i Use `spec()` to retrieve the full column specification for this data.
```

```
i Specify the column types or set `show_col_types = FALSE` to quiet this message.
```

```
x <- cbind(rep(1, nrow(college)), as.numeric(college$out_of_state_tuition)/2000, as.factor(c
```

```
x <- as.matrix(x)
```

```
college <- college |>
```

```
na.omit() |>
```

```
mutate(total = round(out_of_state_total/1000))
```

```
y <- as.vector(college$total)
```

```
trunc_pois_Newton_Raphson(x, y)
```

```
Warning in y - g: longer object length is not a multiple of shorter object  
length
```

```
Warning in y - g: longer object length is not a multiple of shorter object  
length
```

```
Warning in y - g: longer object length is not a multiple of shorter object  
length
```

```
Warning in y - g: longer object length is not a multiple of shorter object  
length
```

```
Warning in y - g: longer object length is not a multiple of shorter object  
length
```

```
Warning in y - g: longer object length is not a multiple of shorter object  
length
```

```
Warning in y - g: longer object length is not a multiple of shorter object  
length
```

```
Warning in y - g: longer object length is not a multiple of shorter object
```


Warning in y - g: longer object length is not a multiple of shorter object length

```
      [,1]
[1,]  3.749873567
[2,]  0.001541439
[3,] -0.015000046
```

Note that this result appears very similar to the result for the zero truncated glm function I found online. Citation ()

Cross Check with Zero Truncated GLM

```
library(VGAM)
```

Loading required package: stats4

Loading required package: splines

```
college <- college |>
  mutate(
    out_tuition = out_of_state_tuition/2000
  )
# Fit a zero-truncated Poisson GLM
model <- vglm(total ~ out_tuition + type, family = pospoisson(), data = college)

summary(model)
```

Call:

```
vglm(formula = total ~ out_tuition + type, family = pospoisson(),
     data = college)
```

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	2.851016	0.027530	103.560	<2e-16 ***
out_tuition	0.052979	0.001289	41.096	<2e-16 ***
typePublic	0.002543	0.018548	0.137	0.891

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Name of linear predictor: `loglink(lambda)`

Log-likelihood: -1580.193 on 545 degrees of freedom

Number of Fisher scoring iterations: 4

No Hauck-Donner effect found in any of the estimates

See, this is extremely close to my output, giving evidence that my Newton Raphson works!

Question 9

By default, `glm()` calls iteratively re-weighted least squares.

Question 10

No, the pareto 1 distribution is not in the exponential family.